

# **The ISRI Analytic Tools**

## **for OCR Evaluation**

**Version 5.1**

**Stephen V. Rice and Thomas A. Nartker**

**Information Science Research Institute**

**TR-96-02**

**August 1996**

Copyright © 1996 by the Information Science Research Institute.

All rights reserved. This document may not be reproduced in whole or in part by any means without permission. For information write: Information Science Research Institute, University of Nevada, Las Vegas, 4505 Maryland Parkway, Box 454021, Las Vegas, Nevada 89154-4021.

# Contents

1	Introduction.....	1
1.1	Operating Environment.....	1
1.2	Special Characters.....	2
1.3	Zoning.....	2
2	Character Accuracy.....	3
2.1	The <b>accuracy</b> Program.....	3
2.2	The <b>synctext</b> Program.....	6
2.3	The <b>accsum</b> Program.....	9
2.4	The <b>groupacc</b> Program.....	12
2.5	The <b>accci</b> Program.....	13
2.6	The <b>accdist</b> Program.....	13
2.7	The <b>ngram</b> Program.....	14
2.8	The <b>vote</b> Program.....	15
3	Word Accuracy.....	16
3.1	The <b>wordacc</b> Program.....	16
3.2	The <b>wordaccsum</b> Program.....	19
3.3	The <b>nonstopacc</b> Program.....	21
3.4	The <b>wordaccci</b> Program.....	21
3.5	The <b>wordaccdist</b> Program.....	22
3.6	The <b>wordfreq</b> Program.....	22
4	Automatic Zoning.....	23
4.1	The <b>editop</b> Program.....	23
4.2	The <b>editopsum</b> Program.....	24
4.3	The <b>editopcost</b> Program.....	25
5	Foreign-Language OCR Testing.....	26
5.1	Latin1 Testing.....	26
5.2	Unicode Testing.....	31
6	Source Files.....	39
	References.....	40



# 1 Introduction

The Information Science Research Institute (ISRI) at the University of Nevada, Las Vegas (UNLV) tested optical character recognition (OCR) systems on an annual basis from 1992 through 1996. The OCR systems that were evaluated are known as “page readers.” These systems take as input a bitmapped image of any document page and attempt to locate and identify the machine-printed characters on the page. Each annual test is described in a technical report (see references [1] through [5]). Synopses of the test results have been published in *Inform* magazine (references [6] through [9]).

Since 1991, ISRI has been actively developing performance measures for page-reading systems. These measures have been used in the annual test and enable a comprehensive evaluation of these systems. The measures include character accuracy, marked character efficiency, word accuracy, non-stopword accuracy, phrase accuracy, and the cost of correcting automatic zoning errors. With the exception of the automatic zoning measure, which is described in [10], a formal specification of these measures, including algorithms for computing them, is presented in Steve Rice’s doctoral dissertation [11].

The ISRI OCR Experimental Environment is a suite of software tools for conducting large-scale, automated tests of page readers. The ISRI Analytic Tools is a subset of this suite and includes the programs that compute the performance measures. These programs are described in this document one by one with examples.

## 1.1 Operating Environment

The Analytic Tools are Unix shell programs that are suitable for use in batch-oriented shell scripts (which are essential to large-scale, automated testing). Each program is given command-line parameters and performs its task in a non-interactive manner. If processing is successful, a zero exit status is returned; otherwise, a non-zero status is returned and an error message is written to “stderr,” the Unix standard error file.

The usage of a program can be displayed by entering the name of the program with no arguments. Specifying “-h” or “-help” as an argument will also display the usage. A *man* page is available for each program.

The programs were developed for Sun SPARCstations under SunOS 4.1.3. They should operate under Solaris without modification. The programs are written in C and could be ported to other platforms.

## 1.2 Special Characters

The user must supply OCR-generated text files and correct (“ground truth”) text files. The Analytic Tools include programs to compare an OCR-generated file with the correct file and obtain measures of performance.

A tilde (~) in an OCR-generated text file is treated as a reject character. A circumflex (^) is interpreted as a suspect marker and serves to mark the following character as suspect. For example, in Ne^vada, the v is marked as suspect. The value of these special characters is assessed when computing marked character efficiency.

Each tilde (~) in a correct text file is treated as a wildcard and allows zero or one arbitrary character to be generated for it without an error being charged. For example, suppose the page contains a character that a page reader is not expected to identify, such as a degree symbol (°) or Greek letter (δ). Then the character should be represented in the correct file by a tilde so that a page reader may generate any single character for it, or no character at all, without penalty. For more information on the use of wildcards in ground truth, see reference [12].

Extraneous spacing characters are ignored in text files and have no effect on the computation of performance measures. Specifically, blank lines are disregarded, as well as leading and trailing blanks on a line. Multiple consecutive blanks within a line are treated as a single blank. The “newline” character at the end of each line is not ignored and appears as “<\n>” within the accuracy reports. Other whitespace characters, such as tabs and formfeeds, are treated as blanks. Page-reading systems and human ground-truth preparers may freely utilize spacing characters to format their text.

## 1.3 Zoning

We assume that an ordered set of zones has been defined for a page image, and that the coordinates of these zones are communicated to the page-reading system when it processes the image. The resulting OCR-generated text file contains the recognized text for zone 1, followed by the text for zone 2, and so on. The correct text file for this page must match this sequence, with the ground truth for zone 1 appearing first, followed by the ground truth for zone 2, and so on. Character and word accuracy are computed from these files as described in Sections 2 and 3. However, for the evaluation of automatic zoning, which is discussed in Section 4, the correspondence between OCR-generated and correct text is not guaranteed. In automatic zoning, the page reader processes the page image without the benefit of zone coordinates; hence, it may produce text in a sequence that does not match the correct file.

## 2 Character Accuracy

### 2.1 The **accuracy** Program

**accuracy** *correctfile* *generatedfile* [ *accuracy\_report* ]

The **accuracy** program compares the correct text found in *correctfile* with the OCR-generated text found in *generatedfile*. A character accuracy report is written to *accuracy\_report* if specified; otherwise, it is written to “stdout,” the Unix standard output file.

We will illustrate this program with an example. The following is a small page image containing two columns of text. Assume that zone 1 contains the left-hand column and zone 2 contains the right-hand column.

crushed under vacuum in stainless steel tubes. Liberated water was extracted at 200°C and converted, using uranium, into hydrogen for D/H analyses. The deuterium content is expressed in parts per thousand difference (per mil) relative to standard mean ocean water (SMOW) [normalized to the V-SMOW/SLAP scale (7)]. The  $\delta D$  values are plotted against age in Fig. 2.

We cannot attribute the changes in deuterium to water-mineral exchange because the water-bearing fractures in the regional carbonate aquifer, feeding the modern (and fossil) flow system, are typically coated with calcite or dolomite (8). This coating precludes the exchange of hydrogen between water and clay minerals during flow from recharge to discharge areas. In fact, the difference in

The following is the correct text for this page. Notice the two wildcards.

crushed under vacuum in stainless steel tubes. Liberated water was extracted at 200°C and converted, using uranium, into hydrogen for D/H analyses. The deuterium content is expressed in parts per thousand difference (per mil) relative to standard mean ocean water (SMOW) [normalized to the V-SMOW/SLAP scale (7)]. The  $\delta D$  values are plotted against age in Fig. 2.

We cannot attribute the changes in deuterium to water-mineral exchange because the water-bearing fractures in the regional carbonate aquifer, feeding the modern (and fossil) flow system, are typically coated with calcite or dolomite (8). This coating precludes the exchange of hydrogen between water and clay minerals during flow from recharge to discharge areas. In fact, the difference in

Here is OCR-generated text for this page. Notice the reject characters and suspect markers.

crushed under vacuum in stainless steel  
tubes. Liberated water was extracted at  
200°C and converted, using uranium,  
into hydrogen for D/H analyses. The  
deuterium content is expressed in parts  
per thousand difference (per mil) relative  
to standard mean ocean water (SMOW)  
[normalized to the V-SMOWISLAP  
scale (7)1. The 6D values are plotted  
against age in Fig. 3.  
We cannot attribute the changes in deu-  
terium to water-mineral exchange be-  
cause the water-bearing ir^acturss in the  
regional carbonaie aquif^er, feeding the  
modern (and fo~sil) now ~vstem. are  
typically coated ~-^ith calci~s or dolomite  
(6). This coating pr-ecludcs the exchanjie  
of hyd^l-ogen bet~^etern water and clay  
minerals during now from ^I^echarge to  
discharge areas. in f,rct, the di~ference in

On the next page is the character accuracy report produced by the **accuracy** program when given these correct and generated files as inputs. A character accuracy report consists of six sections. The first section indicates the number of characters in the ground truth (756), the number of errors made by the page reader (39), and the character accuracy (94.84%). The second section gives the number of reject characters (6), suspect markers (7), and false marks (1), followed by information relating to marked character efficiency: if a user examines the marked characters (1.72% of the text) and corrects the marked errors, the character accuracy will increase (to 96.96%).

Errors are actually edit operations (character insertions, substitutions, and deletions) that are needed to correct the OCR-generated text. The third section of the character accuracy report gives a breakdown of marked errors, unmarked errors, and total errors by edit operation.

The fourth section shows the accuracy by character class. Here the ground-truth characters are divided into classes and the percentage of characters recognized in each class is reported. The total number of ground-truth characters missed (29) is always equal to the number of insertions plus substitutions.

The fifth section lists the “confusions” sorted by the number of errors charged for each. In this example, four errors were charged because n was generated for fl. Since this confusion requires only two edit operations to correct (one insertion and one substitution), then this confusion must have occurred twice to account for a total of four errors.

The sixth and last section of a character accuracy report provides a complete enumeration of the ground-truth characters.



UNLV-ISRI OCR Accuracy Report Version 5.1

756 Characters  
39 Errors  
94.84% Accuracy

6 Reject Characters  
7 Suspect Markers  
1 False Marks  
1.72% Characters Marked  
96.96% Accuracy After Correction

Ins	Subst	Del	Errors	
0	10	6	16	Marked
2	17	4	23	Unmarked
2	27	10	39	Total

Count	Missed	%Right	
117	0	100.00	ASCII Spacing Characters
31	4	87.10	ASCII Special Symbols
6	2	66.67	ASCII Digits
24	1	95.83	ASCII Uppercase Letters
578	22	96.19	ASCII Lowercase Letters
756	29	96.16	Total

Errors	Marked	Correct-Generated
4	0	{fl}-{n}
3	3	{w}-{~.}
2	2	{r}-{I.}
2	2	{r}-{l-}
2	2	{sy}-{~v}
2	2	{te}-{~s}
2	2	{w}-{~.}
2	0	{,}-{.}
2	0	{a}-{,r}
2	0	{e}-{c}
2	0	{e}-{tr}
2	0	{g}-{ji}
1	1	{f}-{~}
1	1	{s}-{~}
1	1	{ }-{.}
1	0	{/}-{I}
1	0	{2}-{3}
1	0	{8}-{6}
1	0	{I}-{i}
1	0	{l}-{1}
1	0	{e}-{s}
1	0	{f}-{i}
1	0	{t}-{i}
1	0	{ }-{ }

Count	Missed	%Right	
20	0	100.00	{<\n>}
97	0	100.00	{ }
5	0	100.00	{( }
5	0	100.00	{ )}
5	2	60.00	{ , }
5	0	100.00	{ - }
7	0	100.00	{ . }
2	1	50.00	{ / }
2	0	100.00	{ 0 }

2	1	50.00	{2}
1	0	100.00	{7}
1	1	0.00	{8}
1	0	100.00	{A}
1	0	100.00	{C}
2	0	100.00	{D}
1	0	100.00	{F}
1	0	100.00	{H}
1	1	0.00	{I}
2	0	100.00	{L}
2	0	100.00	{M}
2	0	100.00	{O}
1	0	100.00	{P}
3	0	100.00	{S}
3	0	100.00	{T}
1	0	100.00	{V}
3	0	100.00	{W}
1	0	100.00	{[}
1	1	0.00	{]}
56	1	98.21	{a}
7	0	100.00	{b}
26	0	100.00	{c}
27	0	100.00	{d}
88	5	94.32	{e}
14	4	71.43	{f}
16	1	93.75	{g}
20	0	100.00	{h}
37	0	100.00	{i}
21	2	90.48	{l}
13	0	100.00	{m}
44	0	100.00	{n}
28	0	100.00	{o}
7	0	100.00	{p}
1	0	100.00	{q}
45	2	95.56	{r}
31	2	93.55	{s}
51	2	96.08	{t}
20	0	100.00	{u}
4	0	100.00	{v}
10	2	80.00	{w}
4	0	100.00	{x}
7	1	85.71	{y}
1	0	100.00	{z}

## 2.2 The **synctext** Program

**synctext** [ -H ] [ -i ] [ -s ] [ -T ] *textfile1 textfile2 ...* > *resultfile*

The **synctext** program can be used to show the alignment of two or more text files. If exactly two input files are specified (one correct file and one OCR-generated file), then the alignment that is computed is the same as the one computed by the **accuracy** program. This allows the user to see where errors occurred within the text. The algorithm used to compute this alignment is described in reference [11].

If more than two input files are specified, or if the “-H” option is given, then the algorithm described in reference [13] is utilized to perform the alignment. This algorithm is also used by the **vote** program to align multiple text files (see Section 2.8).

The “-T” option selects yet another alignment algorithm which can find transposed matches between two input files (see reference [10]). This algorithm is used by the **editop** program for automatic zoning evaluation (see Section 4.1). The “-i” option specifies that the alignment is to be performed on a case insensitive basis (the default is case sensitive), and the “-s” option displays suspect markers in the output.

The following is the output of the **synctext** program when given the example correct and OCR-generated files as inputs. The characters upon which the input files agree are shown first, and everywhere there is a difference, a number in braces identifies a footnote below that indicates what the difference is.

```
=====
crushed under vacuum in stainless steel
tubes. Liberated water was extracted at
200{1}C and converted{2} using uranium,
into hydrogen for D/H analyses. The
deuterium content is expressed in parts
per thousand difference (per mil) relative
to standard mean ocean water (SMOW)
[normalized to the V-SMOW{3}SLAP
scale (7){4}. The {5}D values are plotted
against age in Fig. {6}.
We cannot attribute the changes in d{7}u-
terium to water-mineral exchange be-
cause the water-bearing {8}r{9}actur{10}s in the
regional carbona{11}e aquifer, feeding the
modern (and fo{12}sil) {13}ow {14}stem{15} are
typically coated {16}ith calci{17} or dolomite
({18}). This coating pr{19}eclud{20}s the exchan{21}e
of hyd{22}ogen bet{23}e{24}n water and clay
minerals during {25}ow from {26}echarge to
discharge areas. {27}n f{28}ct, the di{29}ference in

=====
{1}
Correct   {~}
Generated {"}
=====
{2}
Correct   {,}
Generated {..}
=====
{3}
Correct   {/}
Generated {I}
=====
{4}
Correct   {1}
Generated {1}
=====
{5}
Correct   {~}
Generated {6}
```

```

=====
{6}
Correct   {2}
Generated {3}
=====
{7}
Correct   {e}
Generated {c}
=====
{8}
Correct   {f}
Generated {i}
=====
{9}
Correct   {}
Generated {..}
=====
{10}
Correct   {e}
Generated {s}
=====
{11}
Correct   {t}
Generated {i}
=====
{12}
Correct   {s}
Generated {~}
=====
{13}
Correct   {fl}
Generated {n}
=====
{14}
Correct   {sy}
Generated {~v}
=====
{15}
Correct   {,}
Generated {..}
=====
{16}
Correct   {w}
Generated {~-..}
=====
{17}
Correct   {te}
Generated {~s}
=====
{18}
Correct   {8}
Generated {6}
=====
{19}
Correct   {}
Generated {-}
=====
{20}
Correct   {e}
Generated {c}
=====

```

```

{21}
Correct   {g}
Generated {ji}
=====
{22}
Correct   {r}
Generated {l-}
=====
{23}
Correct   {w}
Generated {~.}
=====
{24}
Correct   {e}
Generated {tr}
=====
{25}
Correct   {fl}
Generated {n}
=====
{26}
Correct   {r}
Generated {I.}
=====
{27}
Correct   {I}
Generated {i}
=====
{28}
Correct   {a}
Generated {,r}
=====
{29}
Correct   {f}
Generated {~}
=====

```

## 2.3 The **accsum** Program

```
accsum accuracy_report1 accuracy_report2 ... > accuracy_report
```

The **accsum** program combines two or more character accuracy reports to produce an aggregate report. While the results for a single page are interesting, tabulating the results for a set of pages yields important insights into the page-reading process.

On the next page is an aggregate character accuracy report that was produced by combining 175 individual reports using the **accsum** program. We see that there is a total of 293,493 characters on these 175 pages, and that the page reader made 5,916 errors for an overall character accuracy of 97.98%.

Given this large amount of data, some interesting observations can be made. For example, only 94.53% of the digits were recognized correctly versus 99.07% of the lowercase letters. Only 91.98% of the occurrences of the numeral 1 were correctly

identified.

The list of confusions is very long so it has been truncated here. The most common confusion for this set of pages, contributing 236 errors, is the generation of a space where there should be none, which causes a word to be split (e.g., Nev ada). There are 43 errors attributed to the opposite case, i.e., no space is generated where there should be one, which causes two words to be joined (e.g., LasVegas).

The numeral zero (0) was generated for the letter O a total of 92 times on these pages; it is difficult also for humans to distinguish these symbols. The confusion charged 39 errors was a single sequence of 39 characters that was omitted by the page reader. An ellipsis (...) in a confusion indicates that a long sequence of characters has been truncated.

#### UNLV-ISRI OCR Accuracy Report Version 5.1

293493	Characters			
5916	Errors			
97.98%	Accuracy			
700	Reject Characters			
5	Suspect Markers			
5	False Marks			
0.24%	Characters Marked			
98.37%	Accuracy After Correction			
Ins	Subst	Del	Errors	
71	650	415	1136	Marked
807	2699	1274	4780	Unmarked
878	3349	1689	5916	Total
Count	Missed	%Right		
47850	234	99.51	ASCII Spacing Characters	
14306	603	95.78	ASCII Special Symbols	
14341	784	94.53	ASCII Digits	
18737	768	95.90	ASCII Uppercase Letters	
198259	1838	99.07	ASCII Lowercase Letters	
293493	4227	98.56	Total	
Errors	Marked	Correct-Generated		
236	0	{ }-{ }		
92	0	{0}-{0}		
90	0	{,}-{,}		
61	0	{ }-{ ' }		
52	0	{1}-{1}		
52	0	{ }-{ . }		
43	0	{ }-{ }		
39	0	{tion of<\n>temperature a...}-{ }		
38	0	{.}-{,}		
37	0	{,}-{ }		
31	0	{1}-{I}		
25	0	{-}-{ }		
24	0	{0}-{o}		
24	0	{2}-{z}		
24	0	{in}-{m}		
24	0	{1}-{1}		
21	21	{ }-{~}		
21	0	{1}-{I}		

20	0	{96}-{%}
20	0	{o}-{a}
...		
Count	Missed	%Right
6064	23	99.62 {<\n>}
41786	211	99.50 { }
159	3	98.11 {\"}
13	6	53.85 {#}
3	0	100.00 {\$}
56	13	76.79 {%}
6	1	83.33 {&}
70	3	95.71 {\'}
883	46	94.79 {( }
881	44	95.01 { ) }
75	6	92.00 { * }
54	13	75.93 { + }
3192	150	95.30 { , }
2353	113	95.20 { - }
5620	88	98.43 { . }
224	17	92.41 { / }
2513	110	95.62 { 0 }
2806	225	91.98 { 1 }
1797	122	93.21 { 2 }
1328	41	96.91 { 3 }
1109	57	94.86 { 4 }
1131	54	95.23 { 5 }
1021	65	93.63 { 6 }
858	31	96.39 { 7 }
812	35	95.69 { 8 }
966	44	95.45 { 9 }
234	4	98.29 { : }
157	4	97.45 { ; }
25	0	100.00 { < }
170	38	77.65 { = }
6	2	66.67 { > }
30	2	93.33 { ? }
1267	43	96.61 { A }
512	22	95.70 { B }
1393	31	97.77 { C }
686	29	95.77 { D }
1494	39	97.39 { E }
668	14	97.90 { F }
500	28	94.40 { G }
542	19	96.49 { H }
1097	38	96.54 { I }
198	3	98.48 { J }
177	7	96.05 { K }
710	13	98.17 { L }
774	55	92.89 { M }
1135	45	96.04 { N }
902	147	83.70 { O }
811	38	95.31 { P }
82	7	91.46 { Q }
1187	41	96.55 { R }
1546	41	97.35 { S }
1656	31	98.13 { T }
428	38	91.12 { U }
279	11	96.06 { V }
452	20	95.58 { W }
67	2	97.01 { X }
127	2	98.43 { Y }

47	4	91.49	{z}
33	10	69.70	{[}
33	12	63.64	{]}
3	2	33.33	{^}
17028	111	99.35	{a}
2638	50	98.10	{b}
7269	75	98.97	{c}
7500	32	99.57	{d}
25274	135	99.47	{e}
4819	93	98.07	{f}
3434	50	98.54	{g}
7603	59	99.22	{h}
15387	167	98.91	{i}
135	12	91.11	{j}
851	16	98.12	{k}
9304	156	98.32	{l}
4966	100	97.99	{m}
14430	118	99.18	{n}
15151	107	99.29	{o}
4371	40	99.08	{p}
364	6	98.35	{q}
13574	107	99.21	{r}
12743	96	99.25	{s}
17702	103	99.42	{t}
5627	58	98.97	{u}
2023	22	98.91	{v}
2216	39	98.24	{w}
590	30	94.92	{x}
2898	49	98.31	{y}
362	7	98.07	{z}
3	3	0.00	{{}
20	20	0.00	{ }
3	3	0.00	{}}

## 2.4 The **groupacc** Program

**groupacc** *groupfile* *accuracy\_report* [ *groupacc\_report* ]

In accuracy by character class, ground-truth characters are grouped according to predefined classes, such as the ASCII digits and the ASCII lowercase letters. The **groupacc** program allows user-defined groupings. This program summarizes the accuracy data from *accuracy\_report* for the characters specified in *groupfile*. The results are written to *groupacc\_report* if specified; otherwise, they are written to stdout.

For example, suppose we wish to focus on lowercase letters with descenders. We would create a group file containing the desired letters: *gjpqy*. Given this group file and the aggregate character accuracy report shown above, the **groupacc** program produces the following display:

Count	Missed	%Right	
3434	50	98.54	{g}
135	12	91.11	{j}
4371	40	99.08	{p}



364	6	98.35	{q}
2898	49	98.31	{y}
11202	157	98.60	Total

## 2.5 The **accci** Program

```
accci accuracy_report1 accuracy_report2 ... > resultfile
```

Given a set of character accuracy reports as input, the **accci** program computes an approximate 95% confidence interval for character accuracy using the method of jackknife estimation (which is described in reference [14]). Each input report is treated as one observation. For best results, at least 30 observations are needed.

The following is the output from the **accci** program when given 175 character accuracy reports as input:

```

      175  Observations
293493  Characters
      5916  Errors
      97.98%  Accuracy
97.75%, 98.22%  Approximate 95% Confidence Interval for Accuracy
```

## 2.6 The **accdist** Program

```
accdist accuracy_report1 accuracy_report2 ... > xyfile
```

The **accdist** program computes the distribution of the character accuracies found in a set of character accuracy reports. The results are written in “xy format” to stdout. In this format, each line contains the  $x$ - and  $y$ -coordinates of a data point separated by spaces. A file in this format can be easily imported into most graphing programs.

The **accdist** program produces one data point for each value of  $x$  from 0 to 100. The  $y$  value is the percentage of characters recognized with at least  $x\%$  character accuracy. Below is a portion of the output for the 175-page sample. It shows, for example, that 98.25% of the sample was recognized with at least 95% character accuracy. It is widely believed that it costs more to correct OCR-generated text that is less than 95% accurate than it costs to type the entire text from scratch. Thus, pages recognized with less than 95% accuracy should be sent to a manual data-entry operation. In this example, pages containing 1.75% of the characters should be entered manually.

```

0 100.00
1 100.00
...
94 98.78
```

```

95  98.25
96  95.09
97  88.51
98  63.68
99  16.99
100 0.53

```

## 2.7 The **ngram** Program

**ngram** [-n 1|2|3 ] *textfile1 textfile2 ...* > *resultfile*

The **ngram** program computes  $n$ -gram statistics for one or more text files. If the “-n” option is omitted, or if “-n 1” is specified, this program displays the frequency of each character found in the input files (“unigrams”). If “-n 2” is chosen, the frequency of every distinct character pair is shown (“bigrams”), and if “-n 3” is selected, the frequency of each unique triple of characters (i.e., three consecutive characters) is displayed (“trigrams”).

The  $n$ -gram statistics are displayed twice: first in the collating order of the characters, and then in order by decreasing frequency. Below is abbreviated unigram output for 175 correct text files. The blank character occurs most often, with 41,786 occurrences found in the input files. The letter e is the next most common, with 25,274 occurrences. In this example, none of the occurrences have been marked as suspect.

Count	Suspect	
6064	0	{<\n>}
41786	0	{ }
159	0	{^}
13	0	{#}
3	0	{\$}
56	0	{%}
6	0	{&}
...		
590	0	{x}
2898	0	{y}
362	0	{z}
3	0	{{}
20	0	{ }
3	0	{}}
5453	0	{~}
298946	0	Total

  

Count	Suspect	
41786	0	{ }
25274	0	{e}
17702	0	{t}
17028	0	{a}
15387	0	{i}
15151	0	{o}
14430	0	{n}
...		

13	0	{#}
6	0	{&}
6	0	{>}
3	0	{}
3	0	{`}
3	0	{{}}
3	0	{}}
298946	0	Total

## 2.8 The **vote** Program

**vote** [ -O ] [ -o *outputfile* ] [ -s *m/n* ] [ -w *m/n* ] *textfile1 textfile2 ...*

Given OCR-generated text files that were produced by different page readers for the same page, the **vote** program applies a voting algorithm to produce a more accurate single text file. The input files are first aligned so that agreements and disagreements among the page readers are evident. Then a majority vote is taken to resolve the disagreements. For example, if three page readers believe that a character is an e and two believe it to be a c, then an e will be output by the voting algorithm. (The algorithm is actually more complicated than this, and involves heuristics for breaking ties.) The resulting text is written to *outputfile* if specified; otherwise, it is written to stdout.

The accuracy of the voting output is normally greater than the accuracy of the text from each page reader. ISRI tests have shown that the number of errors in voting output can be as much as 80% less than the number of errors made by the most accurate of the page readers. See references [1], [3], and [5] for voting test results.

The “-O” option enables some important optimizations and should be specified to get the best results. An output character is marked as suspect if it receives no more than the fraction of votes specified by the “-s” option. For example, output characters that receive one-third or less of the possible number of votes will be marked as suspect if “-s 1/3” is specified. The “-w” option indicates the fraction of a vote that each input character receives if it is marked as suspect. This reduces the influence of marked characters on the voting.

## 3 Word Accuracy

### 3.1 The **wordacc** Program

**wordacc** [ *-S stopwordsfile* ] *correctfile generatedfile* [ *wordacc\_report* ]

The **wordacc** program compares the correct text found in *correctfile* with the OCR-generated text found in *generatedfile*. A word accuracy report is written to *wordacc\_report* if specified; otherwise, it is written to stdout.

Only the words found in *stopwordfile* are considered to be stopwords. (If the “-S” option is omitted, the default set of 110 stopwords from the BASISplus text retrieval system is used. See reference [15] for information on this system.) If stopwords are placed in the file in order by decreasing frequency of usage (as determined from some large corpus), then the file can also be used with the **nonstopacc** program (see Section 3.3). Here is an example of such a file containing 200 English stopwords in decreasing order of frequency:

```
the of and to a in that is was he for it with as his on be at by i this had not
are but from or have an they which one you were her all she there would their
we him been has when who will more no if out so said what up its about into
than them can only other new some could these two may then do first any my now
such like our over man me even most made after also did many before must
through back years where much your way well down should because each just those
mr how too state good very make still see men work long get here between both
being under never same another know while last might us great old year off come
since against go came right used take three states himself few use during
without again place around however small mrs thought went say part once general
high upon every does got number until always away something fact though less
put think almost enough far took yet better nothing end why find going asked
later knew point next give group toward young let room side given
```

On the next page is the word accuracy report produced by the **wordacc** program when given this stopwords file and the example correct and generated files. A word accuracy report consists of seven sections. The first section indicates the number of words in the ground truth (119), the number of words misrecognized by the page reader (18), and the word accuracy (84.87%). The second and third sections show the stopwords accuracy (92.86%) and the non-stopword accuracy (80.52%), respectively, with a breakdown by word length for each.

The fourth section presents the results for distinct non-stopword accuracy, which differs from non-stopword accuracy. See reference [4] for a description of this performance measure.

The fifth section gives the phrase accuracy for phrases of lengths 1 through 8. The sixth and seventh sections provide a complete enumeration of the stopwords and non-stopwords, respectively. Word accuracy is determined on a case-insensitive basis, so the

words are displayed here in lowercase only.

UNLV-ISRI OCR Word Accuracy Report Version 5.1

-----  
 119 Words  
 18 Misrecognized  
 84.87% Accuracy

#### Stopwords

Count	Missed	%Right	Length
17	0	100.00	2
16	0	100.00	3
5	2	60.00	4
1	0	100.00	5
1	0	100.00	6
2	1	50.00	7
42	3	92.86	Total

#### Non-stopwords

Count	Missed	%Right	Length
5	0	100.00	1
6	1	83.33	3
7	4	42.86	4
13	0	100.00	5
8	2	75.00	6
11	1	90.91	7
12	3	75.00	8
12	3	75.00	9
3	1	66.67	10
77	15	80.52	Total

#### Distinct Non-stopwords

Count	Missed	%Right	Occurs
58	9	84.48	1
7	1	85.71	2
1	0	100.00	5
66	10	84.85	Total

#### Phrases

Count	Missed	%Right	Length
119	18	84.87	1
118	31	73.73	2
117	39	66.67	3
116	47	59.48	4
115	53	53.91	5
114	57	50.00	6
113	59	47.79	7
112	61	45.54	8

#### Stopwords

Count	Missed	%Right	
1	0	100.00	against
3	0	100.00	and
2	0	100.00	are
1	0	100.00	at
1	0	100.00	be
1	1	0.00	between
1	0	100.00	during
1	1	0.00	fact
1	0	100.00	for
1	0	100.00	from
7	0	100.00	in

1	0	100.00	into
1	0	100.00	is
1	0	100.00	of
1	0	100.00	or
9	0	100.00	the
1	0	100.00	this
4	0	100.00	to
1	0	100.00	under
1	0	100.00	was
1	0	100.00	we
1	1	0.00	with

# Non-stopwords

Count	Missed	%Right	
1	0	100.00	age
1	0	100.00	analyses
1	0	100.00	aquifer
1	0	100.00	areas
1	0	100.00	attribute
1	0	100.00	bearing
1	0	100.00	c
1	1	0.00	calcite
1	0	100.00	cannot
1	1	0.00	carbonate
1	0	100.00	cause
1	0	100.00	changes
1	0	100.00	clay
1	0	100.00	coated
1	0	100.00	coating
1	0	100.00	content
1	0	100.00	converted
1	0	100.00	crushed
2	0	100.00	d
1	1	0.00	deu
1	0	100.00	deuterium
2	1	50.00	difference
1	0	100.00	discharge
1	0	100.00	dolomite
2	1	50.00	exchange
1	0	100.00	expressed
1	0	100.00	extracted
1	0	100.00	feeding
1	0	100.00	fig
2	2	0.00	flow
1	1	0.00	fossil
1	1	0.00	fractures
1	0	100.00	h
2	1	50.00	hydrogen
1	0	100.00	liberated
1	0	100.00	mean
1	0	100.00	mil
1	0	100.00	mineral
1	0	100.00	minerals
1	0	100.00	modern
1	0	100.00	normalized
1	0	100.00	ocean
1	0	100.00	parts
2	0	100.00	per
1	0	100.00	plotted
1	1	0.00	precludes
1	1	0.00	recharge

1	0	100.00	regional
1	0	100.00	relative
1	0	100.00	scale
1	1	0.00	slap
2	1	50.00	smow
1	0	100.00	stainless
1	0	100.00	standard
1	0	100.00	steel
1	1	0.00	system
1	0	100.00	terium
1	0	100.00	thousand
1	0	100.00	tubes
1	0	100.00	typically
1	0	100.00	uranium
1	0	100.00	using
1	0	100.00	v
1	0	100.00	vacuum
1	0	100.00	values
5	0	100.00	water

## 3.2 The **wordaccsum** Program

**wordaccsum** *wordacc\_report1 wordacc\_report2 ... > wordacc\_report*

The **wordaccsum** program combines two or more word accuracy reports to produce an aggregate report. Below is the aggregate report that was produced by combining 175 individual reports. Of 43,928 words on these 175 pages, 2,211 were misrecognized for an overall word accuracy of 94.97%. It is usual to see a higher overall stopword accuracy (97.80%) than non-stopword accuracy (93.22%). The lists of stopwords and non-stopwords are very long and have been truncated here.

UNLV-ISRI OCR Word Accuracy Report Version 5.1

```
-----
43928  Words
2211   Misrecognized
94.97% Accuracy
```

### Stopwords

Count	Missed	%Right	Length
810	46	94.32	1
6069	128	97.89	2
6065	92	98.48	3
2438	61	97.50	4
949	27	97.15	5
204	10	95.10	6
232	5	97.84	7
1	0	100.00	9
16768	369	97.80	Total

### Non-stopwords

Count	Missed	%Right	Length
2275	250	89.01	1
1050	249	76.29	2

1425	169	88.14	3
3025	160	94.71	4
3705	186	94.98	5
3327	172	94.83	6
3216	139	95.68	7
3132	166	94.70	8
2345	125	94.67	9
1560	96	93.85	10
1032	57	94.48	11
599	38	93.66	12
260	17	93.46	13
118	5	95.76	14
58	11	81.03	15
23	2	91.30	16
3	0	100.00	17
5	0	100.00	18
2	0	100.00	19
27160	1842	93.22	Total

#### Distinct Non-stopwords

Count	Missed	%Right	Occurs
11889	808	93.20	1
2437	54	97.78	2
868	15	98.27	3
431	7	98.38	4
241	1	99.59	5
155	4	97.42	6
107	4	96.26	7
69	1	98.55	8
62	1	98.39	9
35	0	100.00	10
112	1	99.11	>10
16406	896	94.54	Total

#### Phrases

Count	Missed	%Right	Length
43928	2211	94.97	1
43753	3927	91.02	2
43578	5446	87.50	3
43403	6812	84.31	4
43228	8082	81.30	5
43053	9249	78.52	6
42878	10334	75.90	7
42704	11351	73.42	8

#### Stopwords

Count	Missed	%Right	
723	32	95.57	a
69	1	98.55	about
32	1	96.88	after
9	1	88.89	again
10	0	100.00	against

...

#### Non-stopwords

Count	Missed	%Right	
2	2	0.00	ab
3	0	100.00	abbreviations
3	0	100.00	abernethy
6	0	100.00	ability
6	0	100.00	able



...

### 3.3 The **nonstopacc** Program

**nonstopacc** *stopwordfile wordacc\_report > xyfile*

Given an ordered list of  $N$  stopwords in *stopwordfile* and a word accuracy report in *wordacc\_report*, the **nonstopacc** program writes the data for the non-stopword accuracy curve to stdout. This curve presents non-stopword accuracy as a function of the number of stopwords. Examples of these curves can be found in reference [5]. The data points are written in *xy* format for values of  $x$  ranging from 0 to  $N$ . Each  $y$  value is the non-stopword accuracy when using the  $x$  most frequently-occurring stopwords from the stopwords file. Here is a portion of the output for the example stopwords file and aggregate word accuracy report:

```
0  94.97
1  94.69
2  94.50
3  94.35
4  94.26
5  94.23
6  94.13
7  94.09
8  94.02
9  94.00
...
200 93.22
```

### 3.4 The **wordaccci** Program

**wordaccci** *wordacc\_report1 wordacc\_report2 ... > resultfile*

Given a set of word accuracy reports as input, the **wordaccci** program computes an approximate 95% confidence interval for word accuracy using the method of jackknife estimation [14]. This program is analogous to the **accci** program for character accuracy (see Section 2.5). The following is the output from the **wordaccci** program when given 175 word accuracy reports as input:

```
175  Observations
43928 Words
2211 Misrecognized
94.97% Accuracy
94.40%, 95.54% Approximate 95% Confidence Interval for Accuracy
```

### 3.5 The **wordaccdist** Program

```
wordaccdist wordacc_report1 wordacc_report2 ... > xyfile
```

The **wordaccdist** program computes the distribution of the word accuracies found in a set of word accuracy reports. The results are written in *xy* format to stdout. This program is analogous to the **accdist** program for character accuracy (see Section 2.6).

The **wordaccdist** program produces one data point for each value of  $x$  from 0 to 100. The  $y$  value is the percentage of words recognized with at least  $x\%$  word accuracy.

### 3.6 The **wordfreq** Program

```
wordfreq textfile1 textfile2 ... > resultfile
```

The **wordfreq** program determines the frequency of words in one or more text files. The frequency data are displayed twice: first in the collating order of the words, and then in order by decreasing frequency. Below is abbreviated **wordfreq** output for 175 correct text files. The word “the” occurs most often, with 2,981 occurrences found in the input files.

```
Count
 723  a
    2  ab
    3  abbreviations
    3  abernethy
    6  ability
...
   28  zone
    9  zones
    2  zonplt
    1  zro
    3  zubovic
43928 Total
```

```
Count
2981  the
1907  of
1366  and
1048  in
 723  a
...
    1  zirconates
    1  zirconium
    1  zirconolite
    1  zonal
    1  zro
43928 Total
```

## 4 Automatic Zoning

### 4.1 The **editop** Program

**editop** *correctfile* *generatedfile* [ *editop\_report* ]

The **editop** program compares the correct text found in *correctfile* with the OCR-generated text found in *generatedfile*. An edit operation report is written to *editop\_report* if specified; otherwise, it is written to stdout. This program is used to evaluate the automatic zoning capability of a page reader.

In automatic zoning, the page reader attempts to locate all text regions on the page and determine their correct reading order. Character recognition is then performed for each text region that has been identified. The OCR-generated text reflects a combination of automatic zoning errors and character recognition errors.

The **editop** program estimates the number of edit operations needed to correct the OCR-generated text. Three types of edit operations are considered: character insertions, character deletions, and block move operations. If a page reader fails to find a text region, insertions are needed to enter the missing text; if it misidentifies a graphic region as text, it may generate extraneous characters that need to be deleted; and if it incorrectly determines the reading order, block move operations are required to re-order the text.

If a page reader performs automatic zoning on the example page and fails to identify the two columns, it may produce output such as this:

```
crushed under vacuum in stainless steel We cannot attribute the changes in dcu-
tubes. Liberated water was extracted at terium to water-mineral exchange be-
200°C and converted. using uranium, cause the water-bearing ir^.actursss in the
into hydrogen for D/H analyses. The regional carbonaie aquif^er, feeding the
deuterium content is expressed in parts modern (and fo~sil) now ~vstem. are
per thousand difference (per mil) relative typically coated ~-^.ith calci~s or dolomite
to standard mean ocean water (SMOW) (8). This coating pr-ecludcs the exchanjie
[normalized to the V-SMOWISLAP of hyd^1-ogen bet~^.etrn water and clay
scale (7)1. The 6D values are plotted minerals during now from ^I^echarge to
against age in Fig. 3. discharge areas. in f,rct, the di~ference in
```

On the next page is the edit operation report produced by the **editop** program when given the correct file and this generated file as inputs. An edit operation report consists of two sections. The first section gives the number of insertions (30), deletions (40), and move operations (21) for correcting the generated text. The second section gives a breakdown of the move operations by length. In this example, there were 11 single-character moves, one move of 22 characters, one move of 30 characters, etc. (The count shown for length 100 is actually the number of moves of length 100 or more.)

```
-----
      30   Insertions
      40   Deletions
      21   Moves
```

## Moves

Count	Length
11	1
1	22
1	30
1	33
1	34
1	37
1	38
1	39
1	69
1	76
1	100

## 4.2 The **editopsum** Program

**editopsum** *editop\_report1 editop\_report2 ... > editop\_report*

The **editopsum** program combines two or more edit operation reports to produce an aggregate report. Below is the aggregate report that was produced by combining 175 individual reports. The distribution of the lengths of the 2,206 moves has been abbreviated here. There are 68 moves of length 100 or more.

```
-----
    2926   Insertions
   28787   Deletions
    2206   Moves
```

## Moves

Count	Length
1543	1
204	2
97	3
49	4
28	5
25	6
9	7
7	8
7	9
6	10
...	
1	77
1	79
1	84
2	96
68	100

### 4.3 The **editopcost** Program

```
editopcost editop_report [ editop_report2 ] > xyfile
```

The **editopcost** program computes the cost of the edit operations described in *editop\_report*, less the cost of the edit operations described in *editop\_report2*, if specified. The cost is based on the number of insertions, the number and lengths of move operations, and a threshold value used to convert move operations into an equivalent number of insertions. See reference [10] for details. The output is written to stdout in *xy* format. There is a data point specifying a threshold value (*x*) and the associated cost (*y*) for each threshold in the range 0 to 100.

Normally, *editop\_report* is an aggregate report that indicates the edit operations for correcting OCR-generated text that was produced with automatic zoning. The cost of these operations is the total cost of correcting the automatic zoning errors and the character recognition errors. Hence, it is often desirable to specify a second aggregate report, *editop\_report2*, that indicates the edit operations for correcting the OCR-generated text that was produced using manually-defined zones for the same set of pages. The cost of these operations is the cost of correcting only the character recognition errors. Deducting this cost from the total yields the cost of correcting the automatic zoning errors.

Below is an abbreviated example of **editopcost** output. When plotted, these data points form a curve showing the cost of correcting automatic zoning errors as a function of the threshold value. Examples of such curves can be seen in references [3], [4], and [5]. The costs given in the **editopcost** output are unnormalized but can be normalized by dividing each *y* value by the total number of ground-truth characters.

0	77
1	1010
2	1459
3	1803
4	2088
...	
96	11710
97	11767
98	11824
99	11881
100	11938

## 5 Foreign-Language OCR Testing

### 5.1 Latin1 Testing

The examples in the preceding sections illustrate how to use the ISRI Analytic Tools for evaluating English OCR using ASCII text files. These tools can also be used with Latin1 text files to evaluate OCR systems for any of the following languages: Danish, Dutch, Faroese, Finnish, French, German, Icelandic, Irish, Italian, Norwegian, Portuguese, Spanish, and Swedish. Latin1 is the ISO 8859-1 standard 8-bit character encoding [16]. This encoding contains the 7-bit ASCII standard as a subset.

In this section, we illustrate how the Analytic Tools can be used to evaluate Spanish OCR. Here is a small page image containing Spanish text:

**Con la incorporación de técnicos con bajo perfil político y la presencia de hombres con acceso directo a los principales despachos de la Casa Rosada, la conformación del nuevo gabinete municipal parece haber fortalecido la figura del intendente porteño, Saúl Bouer.**

Latin1 is needed to represent the correct text because of the accented characters:

Con la incorporación de técnicos  
con bajo perfil político y la presen-  
cia de hombres con acceso directo a  
los principales despachos de la  
Casa Rosada, la conformación del  
nuevo gabinete municipal parece  
haber fortalecido la figura del in-  
tendente porteño, Saúl Bouer.

Here is OCR-generated text for this page, also in Latin1:

Con la incorporación de técnicos  
con bajo perfil político y la pre^sen-  
cia de hombres con acceso directo a  
los principales despachos de la  
Cas^a Rosada, la conformación del  
nuevo gabinete municipal parece  
haber fortalecido la figura del i^ii^.  
tendente porteño, Sañí Boner.

On the next page is the character accuracy report produced by the **accuracy** program when given these correct and generated files as inputs.

UNLV-ISRI OCR Accuracy Report Version 5.1

270 Characters  
7 Errors  
97.41% Accuracy

1 Reject Characters  
4 Suspect Markers  
2 False Marks  
1.85% Characters Marked  
98.89% Accuracy After Correction

Ins	Subst	Del	Errors	
0	3	1	4	Marked
0	3	0	3	Unmarked
0	6	1	7	Total

Count	Missed	%Right	
43	0	100.00	ASCII Spacing Characters
5	2	60.00	ASCII Special Symbols
5	0	100.00	ASCII Uppercase Letters
211	3	98.58	ASCII Lowercase Letters
6	1	83.33	Latin1 Lowercase Letters
270	6	97.78	Total

Errors	Marked	Correct-Generated
3	3	{n-}-{ii.}
2	0	{ûl}-{ñi}
1	1	{-}-{~}
1	0	{u}-{n}

Count	Missed	%Right	
8	0	100.00	{<\n>}
35	0	100.00	{ }
2	0	100.00	{,}
2	2	0.00	{-}
1	0	100.00	{.}
1	0	100.00	{B}
2	0	100.00	{C}
1	0	100.00	{R}
1	0	100.00	{S}
24	0	100.00	{a}
4	0	100.00	{b}
18	0	100.00	{c}
10	0	100.00	{d}
25	0	100.00	{e}
4	0	100.00	{f}
2	0	100.00	{g}
3	0	100.00	{h}
16	0	100.00	{i}
1	0	100.00	{j}
14	1	92.86	{l}
3	0	100.00	{m}
16	1	93.75	{n}
23	0	100.00	{o}
10	0	100.00	{p}
14	0	100.00	{r}
10	0	100.00	{s}
8	0	100.00	{t}
4	1	75.00	{u}
1	0	100.00	{v}

1	0	100.00	{y}
1	0	100.00	{é}
1	0	100.00	{í}
1	0	100.00	{ñ}
2	0	100.00	{ó}
1	1	0.00	{ú}

The following is the output from the **synctext** program when given these correct and generated files as inputs:

```
=====
Con la incorporación de técnicos
con bajo perfil político y la presen{1}
cia de hombres con acceso directo a
los principales despachos de la
Casa Rosada, la conformación del
nuevo gabinete municipal parece
haber fortalecido la figura del i{2}
tendente porteño, Sa{3} Bo{4}er.

=====
{1}
Correct   {-}
Generated {~}
=====
{2}
Correct   {n-}
Generated {ii.}
=====
{3}
Correct   {úl}
Generated {ñí}
=====
{4}
Correct   {u}
Generated {n}
=====
```

The **wordacc** program needs a stopword file. Here is a file containing 200 Spanish stopwords in decreasing order of frequency (courtesy of Chris Buckley at Cornell):

de la el y en que a los del se por un con las para al una su no es lo como más  
sus pero dijo este ya fue esta entre ha también dos son o está sin le sobre si  
ser cuando hasta porque tiene donde desde parte sólo han todo muy hoy durante  
hay tres quien están uno así todos después además otros expresó hace nuevo  
ahora agregó primera hacer ante señaló les ese e será puede vez ayer mismo  
tienen fueron cada contra aunque pasado mayor lugar otro antes nos mientras esa  
esto ellos algunos primer gran tanto sido otra indicó nuevos eso bien menos  
estos cuatro explicó embargo tener ni debe otras mejor había momento cual  
informó era mucho luego hecho sino nueva pues sea quienes dentro qué cuenta  
cinco me va según unos manera él comentó dar nada muchos sí aún pueden estar  
siempre poco todas haber aquí tan segundo hizo ver toda fin yo casi podría  
estas hacia seis algunas próximo aseguró decir bajo fuera varios misma  
cualquier total estamos algo nosotros añadió mi grandes estaba ello través dio  
ex afirmó tal tenemos existe últimos conocer respecto sería van dice primero  
segunda cosas actualmente



Below is the word accuracy report produced by the **wordacc** program when given this stopword file, the correct file, and the OCR-generated file:

UNLV-ISRI OCR Word Accuracy Report Version 5.1

-----  
 43 Words  
 3 Misrecognized  
 93.02% Accuracy

Stopwords

Count	Missed	%Right	Length
2	0	100.00	1
8	0	100.00	2
6	0	100.00	3
1	0	100.00	4
2	0	100.00	5
19	0	100.00	Total

Non-stopwords

Count	Missed	%Right	Length
1	1	0.00	2
1	0	100.00	3
2	1	50.00	4
1	1	0.00	5
6	0	100.00	6
3	0	100.00	7
4	0	100.00	8
2	0	100.00	9
2	0	100.00	11
1	0	100.00	12
1	0	100.00	13
24	3	87.50	Total

Distinct Non-stopwords

Count	Missed	%Right	Occurs
24	3	87.50	1
24	3	87.50	Total

Phrases

Count	Missed	%Right	Length
43	3	93.02	1
42	4	90.48	2
41	5	87.80	3
40	5	87.50	4
39	5	87.18	5
38	5	86.84	6
37	5	86.49	7
36	5	86.11	8

# Stopwords

Count	Missed	%Right	
1	0	100.00	a
1	0	100.00	bajo
3	0	100.00	con
3	0	100.00	de
2	0	100.00	del
1	0	100.00	haber
5	0	100.00	la
1	0	100.00	los
1	0	100.00	nuevo
1	0	100.00	y

# Non-stopwords

Count	Missed	%Right	
1	0	100.00	acceso
1	1	0.00	bouer
1	0	100.00	casa
1	0	100.00	cia
1	0	100.00	conformación
1	0	100.00	despachos
1	0	100.00	directo
1	0	100.00	figura
1	0	100.00	fortalecido
1	0	100.00	gabinete
1	0	100.00	hombres
1	1	0.00	in
1	0	100.00	incorporación
1	0	100.00	municipal
1	0	100.00	parece
1	0	100.00	perfil
1	0	100.00	político
1	0	100.00	porteño
1	0	100.00	presen
1	0	100.00	principales
1	0	100.00	rosada
1	1	0.00	saúl
1	0	100.00	tendente
1	0	100.00	técnicos

## 5.2 Unicode Testing

The Unicode standard specifies a 16-bit character encoding for nearly all of the world's languages [17]. It includes Latin1 and ASCII as subsets. By using the Unicode representation of characters, the ISRI Analytic Tools can be utilized to evaluate OCR systems for almost any language in the world.

Each program of the Analytic Tools operates on files in “Extended ASCII” format. In this format, ASCII and Latin1 characters are stored directly in the file, and each 16-bit Unicode symbol is represented by four hexadecimal digits surrounded by angle brackets. For example, the sequence **AÄΩℵ** is stored as **AÄ<03A9><05D0>** because **A** is an ASCII character, **Ä** is a Latin1 character, and **Ω** and **ℵ** are represented by codes 03A9 and 05D0 respectively in Unicode.

While all operations are performed using Extended ASCII files, the **asc2uni** and **uni2asc** programs are available to convert to and from the standard Unicode file format.

<b>asc2uni</b> < <i>ASCII_file</i> > <i>Unicode_file</i> <b>uni2asc</b> < <i>Unicode_file</i> > <i>ASCII_file</i>
--

For example, the **accuracy** program accepts a correct file and an OCR-generated file in Extended ASCII format and produces a character accuracy report, also in Extended ASCII format. If the correct text was entered using a Unicode editor into a standard Unicode file, then an equivalent Extended ASCII version of this text can be created using the **uni2asc** program. Similarly, if an OCR system produces a standard Unicode file as output, then this generated text can be converted to Extended ASCII using the **uni2asc** program. The **asc2uni** program allows the character accuracy report to be converted to the standard Unicode file format for viewing with a Unicode display tool.

All of the programs presented in Sections 2 and 4 can be used with Extended ASCII files and work in the same way as described in those sections. However, the programs in Section 3 define a word to be any sequence of one or more ASCII or Latin1 letters. Hence, for languages that compose words using other symbols, word accuracy is unavailable.

We illustrate Unicode testing with an example in Japanese. At the top of the next page is a small page image containing Japanese text.

イメージスキャナで読み取った日  
本語活字文書を、そのまま文字デ  
ータに高速変換。今までのパソコ  
ンOCRでは見られない高機能  
を、ソフトウェアと拡張スロット用  
ボードで実現。お手持ちのパソコ  
ンで、膨大なデータの入力を迅速  
に処理でき、文書入力のためのコ  
ストと負担を大幅に軽減します。

Here is the correct text for this page in Extended ASCII:

```
<30A4><30E1><30FC><30B8><30B9><30AD><30E3><30CA><3067><8AAD><307F><53D6><3063><305F><65E5>  
<672C><8A9E><6D3B><5B57><6587><66F8><3092><3001><305D><306E><307E><307E><6587><5B57><30C7>  
<30FC><30BF><306B><9AD8><901F><5909><63DB><3002><4ECA><307E><3067><306E><30D1><30BD><30B3>  
<30F3>OCR<3067><306F><898B><3089><308C><306A><3044><9AD8><6A5F><80FD>  
<3092><3001><30BD><30D5><30C8><30A6><30A7><30A2><3068><62E1><5F35><30B9><30ED><30C3><30C8><7528>  
<30DC><30FC><30C9><3067><5B9F><73FE><3002><304A><624B><6301><3061><306E><30D1><30BD><30B3>  
<30F3><3067><3001><81A8><5927><306A><30C7><30FC><30BF><306E><5165><529B><3092><8FC5><901F>  
<306B><51E6><7406><3067><304D><3001><6587><66F8><5165><529B><306E><305F><3081><306E><30B3>  
<30B9><30C8><3068><8CA0><62C5><3092><5927><5E45><306B><8EFD><6E1B><3057><307E><3059><3002>
```

Here is a Unicode display of this correct text:

イメージスキャナで読み取った日  
本語活字文書を、そのまま文字デ  
ータに高速変換。今までのパソコ  
ンOCRでは見られない高機能  
を、ソフトウェアと拡張スロット用  
ボードで実現。お手持ちのパソコ  
ンで、膨大なデータの入力を迅速  
に処理でき、文書入力のためのコ  
ストと負担を大幅に軽減します。

Here is a Unicode display of OCR-generated text for this page:

イメージスキャナで読み取った日  
本語活字文書在、そのまま文字デ  
ータに二高速変換。含までの／く、ノコ  
ン・c円では見られなし高機能  
在、ソフトウェアと拡張スロット用  
ボードで実現。お手持ちの／く、ノコ  
ンで、一大なデ・タの入力にB速  
に二処理でと文書入力のためのコ  
ストと負担に大幅に軽減します。

The next three pages show a Unicode display of the character accuracy report produced by the **accuracy** program for this correct and generated text.

UNLV-ISRI OCR Accuracy Report Version 5.1

144 Characters  
52 Errors  
63.89% Accuracy

0 Reject Characters  
0 Suspect Markers  
0 False Marks  
0.00% Characters Marked  
63.89% Accuracy After Correction

Ins	Subst	Del	Errors	
0	0	0	0	Marked
1	41	10	52	Unmarked
1	41	10	52	Total

Count	Missed	%Right	
9	0	100.00	ASCII Spacing Characters
3	3	0.00	ASCII Uppercase Letters
7	1	85.71	CJK Symbols and Punctuation
42	8	80.95	Hiragana
38	22	42.11	Katakana
45	8	82.22	CJK Unified Ideographs
144	42	70.83	Total

Errors	Marked	Correct-Generated
4	0	{バソ}-{/く、ノ}
4	0	{バソ}-{/ (、ノ}
4	0	{一タに}-{-タ   二}
3	0	{を}-{在}
3	0	{ロット}-{ロツト}
3	0	{OCR}-{* c 円}
2	0	{い}-{L l}
2	0	{き、}-{と}
2	0	{に}-{  二}
2	0	{を迅}-{在B}
2	0	{エア}-{ウフ}
2	0	{フト}-{ワト}
2	0	{ージ}-{-シ}
2	0	{一タ}-{-タ}
2	0	{一ド}-{-ト}
2	0	{今}-{ 舍}
2	0	{}-{  }
1	0	{デ}-{テ}
1	0	{ト}-{ト}
1	0	{ヤ}-{ヤ}
1	0	{幅}-{幅}
1	0	{書}-{書}
1	0	{滅}-{高}
1	0	{膨}-{一}
1	0	{負}-{員}
1	0	{速}-{連}

Count	Missed	%Right	
9	0	100.00	{<\n>}
1	1	0.00	{C}
1	1	0.00	{O}
1	1	0.00	{R}
4	1	75.00	{、}
3	0	100.00	{。}
1	1	0.00	{い}
1	0	100.00	{お}
1	1	0.00	{き}
1	0	100.00	{し}
1	0	100.00	{す}
1	0	100.00	{そ}
2	0	100.00	{た}
1	0	100.00	{ち}
1	0	100.00	{っ}
6	0	100.00	{で}
2	0	100.00	{と}
2	0	100.00	{な}
3	2	33.33	{に}
6	0	100.00	{の}
1	0	100.00	{は}
4	0	100.00	{ま}
1	0	100.00	{み}
1	0	100.00	{め}
1	0	100.00	{ら}
1	0	100.00	{れ}
4	4	0.00	{を}
1	1	0.00	{ア}
1	0	100.00	{イ}
1	0	100.00	{ウ}
1	1	0.00	{エ}
1	0	100.00	{キ}
3	0	100.00	{コ}
1	1	0.00	{ジ}
3	0	100.00	{ス}
3	2	33.33	{ソ}
2	2	0.00	{タ}
1	1	0.00	{ッ}
2	1	50.00	{デ}
3	3	0.00	{ト}
1	1	0.00	{ド}
1	0	100.00	{ナ}
2	2	0.00	{バ}

1	1	0.00	{フ}
1	0	100.00	{ボ}
1	0	100.00	{メ}
1	1	0.00	{ヤ}
1	1	0.00	{ロ}
2	0	100.00	{ン}
4	4	0.00	{一}
1	1	0.00	{今}
2	0	100.00	{入}
1	0	100.00	{処}
2	0	100.00	{力}
1	0	100.00	{取}
1	0	100.00	{変}
2	0	100.00	{大}
2	0	100.00	{字}
1	0	100.00	{実}
1	1	0.00	{幅}
1	0	100.00	{張}
1	0	100.00	{手}
1	0	100.00	{担}
1	0	100.00	{鉋}
1	0	100.00	{持}
1	0	100.00	{換}
3	0	100.00	{文}
1	0	100.00	{日}
2	1	50.00	{書}
1	0	100.00	{本}
1	0	100.00	{機}
1	0	100.00	{活}
1	1	0.00	{減}
1	0	100.00	{現}
1	0	100.00	{理}
1	0	100.00	{用}
1	0	100.00	{能}
1	1	0.00	{膨}
1	0	100.00	{見}
1	0	100.00	{語}
1	0	100.00	{読}
1	1	0.00	{負}
1	0	100.00	{軽}
1	1	0.00	{迅}
2	1	50.00	{速}
2	0	100.00	{高}

Here is the output from the **synctext** program for this correct and generated text:

```
=====
イメ {1} スキ {2} ナで読み取った日
本語活字文書 {3} 、そのまま文字 {4}
{5} 高 {6} 変換。 {7} までの {8} コ
ン {9} では見られな {10} 高機能
{11} 、 {12} ソ {13} ウ {14} と拡張ス {15} 用
ボ {16} で実現。 {17} お手持ちの {18} コ
ンで、 {19} 大ナデ {20} の入力 {21} 速
{22} 処理で {23} 文 {24} 入力のためのコ
ス {25} と {26} 担 {27} 大 {28} に軽 {29} します。
```

```
=====
{1}
Correct  {ージ}
Generated {ーシ}
```

```
=====
{2}
Correct  {ヤ}
Generated {ヤ}
```

```
=====
{3}
Correct  {を}
Generated {在}
```

```
=====
{4}
Correct  {デ}
Generated {テ}
```

```
=====
{5}
Correct  {ータに}
Generated {ータ | 二}
```

```
=====
{6}
Correct  {速}
Generated {連}
```

```
=====
{7}
Correct  {今}
Generated { 含}
```



{8}  
Correct {パソ}  
Generated {／ (、ノ)}

---

---

{9}  
Correct {OCR}  
Generated {° c 円}

---

---

{10}  
Correct {い}  
Generated {L l}

---

---

{11}  
Correct {を}  
Generated {在}

---

---

{12}  
Correct {}  
Generated { }

---

---

{13}  
Correct {フト}  
Generated {ワト}

---

---

{14}  
Correct {エア}  
Generated {ウフ}

---

---

{15}  
Correct {ロット}  
Generated {口ツト}

---

---

{16}  
Correct {ード}  
Generated {ート}

---

---

{17}  
Correct {}  
Generated { }

---

---

{18}  
Correct {パソ}  
Generated {／<、ノ}

---

---

{19}  
Correct {膨}  
Generated {一}

---

---

{20}  
Correct {一タ}  
Generated {-タ}

---

---

{21}  
Correct {を迅}  
Generated {在B}

---

---

{22}  
Correct {に}  
Generated {一二}

---

---

{23}  
Correct {き、}  
Generated {と}

---

---

{24}  
Correct {書}  
Generated {善}

---

---

{25}  
Correct {ト}  
Generated {ト}

---

---

{26}  
Correct {負}  
Generated {員}

---

---

{27}  
Correct {を}  
Generated {在}

---

---

{28}  
Correct {幅}  
Generated {福}

---

---

{29}  
Correct {減}  
Generated {高}

---

---

## 6 Source Files

Version 5.1 of the ISRI Analytic Tools consists of the 19 programs described in this report. There is one C source file for each program:

accci.c	editop.c	nonstopacc.c	wordaccci.c
accdist.c	editopcost.c	synctext.c	wordaccdist.c
accsum.c	editopsum.c	uni2asc.c	wordaccsum.c
accuracy.c	groupacc.c	vote.c	wordfreq.c
asc2uni.c	ngram.c	wordacc.c	

Each program is linked with a library of 15 modules. There is a “.c” file and a “.h” include file for each module. For example, the accrpt module consists of the files accrpt.c and accrpt.h. Here is a brief description of these modules:

accrpt	reading and writing character accuracy reports
charclass	determining character classes
ci	computing confidence intervals
dist	computing the distribution of accuracies
edorpt	reading and writing edit operation reports
list	linked list operations
sort	sorting
stopword	identifying stopwords
sync	aligning strings
table	hash table operations
text	reading and writing text files
unicode	reading and writing Unicode files
util	utility routines
wacrpt	reading and writing word accuracy reports
word	parsing to extract words

## References

- [1] S. V. Rice, J. Kanai, and T. A. Nartker. A report on the accuracy of OCR devices. Technical Report 92-02, Information Science Research Institute, University of Nevada, Las Vegas, March 1992.
- [2] S. V. Rice, J. Kanai, and T. A. Nartker. An evaluation of OCR accuracy. Technical Report 93-01, Information Science Research Institute, University of Nevada, Las Vegas, April 1993.
- [3] S. V. Rice, J. Kanai, and T. A. Nartker. The third annual test of OCR accuracy. Technical Report 94-03, Information Science Research Institute, University of Nevada, Las Vegas, April 1994.
- [4] S. V. Rice, F. R. Jenkins, and T. A. Nartker. The fourth annual test of OCR accuracy. Technical Report 95-04, Information Science Research Institute, University of Nevada, Las Vegas, April 1995.
- [5] S. V. Rice, F. R. Jenkins, and T. A. Nartker. The fifth annual test of OCR accuracy. Technical Report 96-01, Information Science Research Institute, University of Nevada, Las Vegas, April 1996.
- [6] T. A. Nartker, S. V. Rice, and J. Kanai. OCR accuracy: UNLV's second annual test. *Inform*, Association for Information and Image Management, 8(1):40+, January 1994.
- [7] T. A. Nartker and S. V. Rice. OCR accuracy: UNLV's third annual test. *Inform*, Association for Information and Image Management, 8(8):30+, September 1994.
- [8] T. A. Nartker, S. V. Rice, and F. R. Jenkins. OCR accuracy: UNLV's fourth annual test. *Inform*, Association for Information and Image Management, 9(7):38+, July 1995.
- [9] T. A. Nartker, S. V. Rice, and F. R. Jenkins. OCR accuracy: UNLV's fifth annual test. *Inform*, Association for Information and Image Management, to appear, 1996.
- [10] J. Kanai, S. V. Rice, T. A. Nartker, and G. Nagy. Automated evaluation of OCR zoning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 17(1):86-90, 1995.
- [11] S. V. Rice. *Measuring the Accuracy of Page-Reading Systems*. Ph.D. dissertation,

University of Nevada, Las Vegas, 1996.

- [12] S. V. Rice, J. Kanai, and T. A. Nartker. Preparing OCR test data. Technical Report 93-08, Information Science Research Institute, University of Nevada, Las Vegas, June 1993.
- [13] S. V. Rice, J. Kanai, and T. A. Nartker. An algorithm for matching OCR-generated text strings. *International Journal of Pattern Recognition and Artificial Intelligence*, 8(5):1259-1268, 1994.
- [14] E. J. Dudewicz and S. N. Mishra. *Modern Mathematical Statistics*, pages 743-748. John Wiley & Sons, 1988.
- [15] Information Dimensions, Inc., Dublin, Ohio. *BASISplus Database Administration Reference, Release L*, June 1990.
- [16] International Organization for Standardization, Geneva, Switzerland. *Information Processing — 8-bit Single-byte Coded Graphic Character Sets: Part 1. Latin Alphabet No. 1*, 1987.
- [17] The Unicode Consortium, San Jose, California. *The Unicode Standard: Worldwide Character Encoding Version 1.0*, Addison-Wesley, Volume 1, 1991, and Volume 2, 1992.